

# Kubernetes Monitoring

## Agenda

1. Vorbereitung
  - [Self-Service Cluster ausrollen](#)
2. Kubernetes Monitoring Grundlagen
  - [Abgrenzung zu Observeability](#)
3. Kubernetes Monitoring (Single Cluster / Instance of Prometheus)
  - [Prometheus Monitoring Server \(Overview\)](#)
  - [Prometheus - Achtung bitte kein prometheus agent](#)
  - [Prometheus / Grafana Stack installieren \(advanced\)](#)
  - [Prometheus / blackbox exporter](#)
4. Prometheus Praxis
  - [Nginx mit ServiceMonitor und export konfigurieren \(sidecar\)](#)
5. Grafana - Dashboards
  - [Bestehendes Dashboard anpassen](#)
  - [Pod und Container Dashboard](#)
6. Grafana - Alerting and Notifications
  - [Grafana neuen alert anlegen](#)
  - [Grafana absence alert konfigurieren - d.h. Service hat keine Pods mehr](#)
  - [Grafana alert, >= pod aus replicaset nicht erreichbar](#)
  - [Grafana Notifications/Contact points](#)
7. Kubernetes Multi-Cluster (Types of setups including disadvantages/advantages)
  - [Recommended: Variant 1: prometheus agent + thanos/grafana stack](#)
  - [Variant 2: Full prometheus in each cluster with thanos sidecar](#)
8. Grafana Loki
  - [Installation von Grafana Loki - Single Instance - für Testing](#)
  - [Datasource in Grafana bereitstellen per helm](#)
  - [Wo finde ich Loki in Grafana ?](#)

## Backlog / Sammlung

1. Prometheus
  - [Prometheus-Metriktypen \(engl. metric types\)](#)
2. Kubernetes Multi-Cluster (using Thanos)
  - [Prerequisites: What is Thanos](#)
  - [Components](#)
  - [Thanos Compactor](#)
3. Kubernetes Multi-Cluster (using Cortex - multi-tenant tsdb's)

## Vorbereitung

### Self-Service Cluster ausrollen

- ausgerollt mit terraform (binary ist installiert) - snap install --classic terraform
- beinhaltet
  1. 1 controlplane
  2. 2 worker nodes
  3. metallb mit ips der Nodes (hacky but works)
  4. ingress mit wildcard-domain: \*.tlx.do.t3isp.de

### Walkthrough

```
cd
git clone https://github.com/jmetzger/training-kubernetes-monitoring-stack-do-terraform.git
cd install
cat /tmp/.env
source /tmp/.env
terraform init
terraform apply -auto-approve
```

### Hinweis

```
## Sollte es nicht sauber durchlaufen
## einfach nochmal
terraform apply -auto-approve

## Wenn das nicht geht, einfach nochmal neu
terraform destroy -auto-approve
terraform apply -auto-approve
```

## Kubernetes Monitoring Grundlagen

### Abgrenzung zu Observeability

Der Unterschied zwischen **Kubernetes Monitoring** und **Observability** liegt vor allem im **Ziel**, **Umfang** und **Vorgehen**. Hier eine klare Gegenüberstellung:

---

#### Monitoring (Beobachtung)

**Definition:** Monitoring ist das **Sammeln und Anzeigen** von Metriken und Logs, um den **Zustand eines Systems zu überwachen**.

#### Merkmale:

- Fokus auf bekannte Metriken, Fehler und Schwellenwerte
- Alarme bei bekannten Problemen (z. B. CPU-Auslastung > 90 %)
- Reaktiv: Man erkennt, *dass* etwas nicht stimmt

#### Beispiele:

- CPU-, RAM-, Netzwerkverbrauch von Pods
- Anzahl von HTTP 500 Errors
- Alerts bei Pod-Crashes

#### Werkzeuge:

- Prometheus
  - Grafana
  - Alertmanager
  - Metrics Server
-

## Observability (Beobachtbarkeit)

**Definition:** Observability beschreibt die **Fähigkeit**, den **internen Zustand** eines Systems allein durch seine externen Outputs (Logs, Metriken, Traces) **verstehen** zu können.

### Merkmale:

- Ermöglicht Ursachenanalyse auch für **unbekannte Probleme**
- Proaktiv: Man kann *warum* etwas passiert ist, herausfinden
- Nutzt drei Hauptsäulen:
  - **Metriken** (z. B. Requests/sec)
  - **Logs** (z. B. Stacktraces)
  - **Traces** (z. B. verteilte Aufrufe zwischen Microservices)

### Beispiele:

- Warum ist ein Request langsam? (Trace-Analyse)
- Welche Kette von Services war beteiligt?
- Korrelation von Logs mit Metrik-Anomalien

### Werkzeuge:

- OpenTelemetry
- Grafana Tempo / Jaeger (für Tracing)
- Loki (für Logs)
- Elastic Stack
- Honeycomb, Lightstep

---

### ⚡ Vergleich zusammengefasst

Aspekt	Monitoring	Observability
Ziel	Systemzustand überwachen	Systemverhalten verstehen
Fokus	Bekannte Probleme erkennen	Ursachen auch unbekannter Probleme analysieren
Methoden	Metriken, Schwellenwerte, Alarmer	Logs, Metriken, Traces kombiniert
Reaktiv / Proaktiv	Reaktiv	Proaktiv / Diagnostisch
Typische Fragen	Ist etwas kaputt?	Warum ist etwas kaputt?

---

## Kubernetes Monitoring (Single Cluster / Instance of Prometheus)

### Prometheus Monitoring Server (Overview)

#### What does it do ?

- It monitors your system by collecting data
- Data is pulled from your system by defined endpoints (http) from your cluster
- To provide data on your system, a lot of exporters are available, that
  - collect the data and provide it in Prometheus

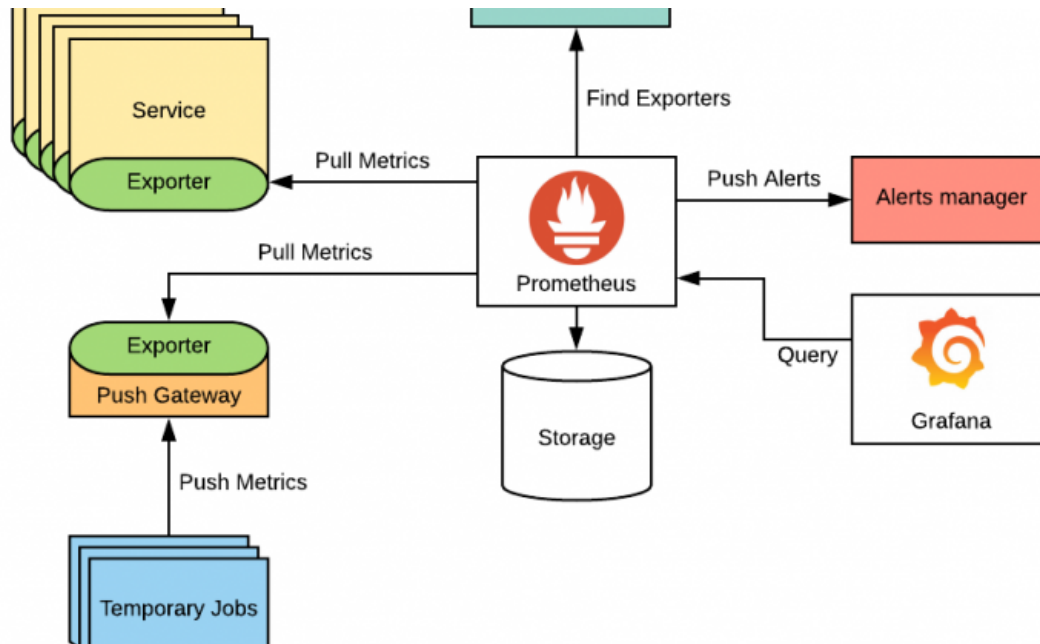
#### Technical

- Prometheus has a TDB (Time Series Database) and is good as storing time series with data
- Prometheus includes a local on-disk time series database, but also optionally integrates with remote storage systems.
- Prometheus's local time series database stores data in a custom, highly efficient format on local storage.
- Ref: <https://prometheus.io/docs/prometheus/latest/storage/>

#### What are time series ?

- A time series is a sequence of data points that occur in successive order over some period of time.
- Beispiel:
  - Du willst die täglichen Schlusspreise für eine Aktie für ein Jahr dokumentieren
  - Damit willst Du weitere Analysen machen
  - Du würdest das Paar Datum/Preis dann in der Datumsreihenfolge sortieren und so ausgeben
  - Dies wäre eine "time series"

## Komponenten von Prometheus



Quelle: <https://www.devopsschool.com/>

## Prometheus Server

1. Retrieval (Sammeln)
  - Data Retrieval Worker
    - pull metrics data
2. Storage
  - Time Series Database (TDB)
    - stores metrics data
3. HTTP Server
  - Accepts PromQL - Queries (e.g. from Grafana)
    - accept queries

## Grafana ?

- Grafana wird meist verwendet um die grafische Auswertung zu machen.
- Mit Grafana kann ich einfach Dashboards verwenden
- Ich kann sehr leicht festlegen (Durch Data Sources), wo meine Daten herkommen

## Prometheus - Achtung bitte kein prometheus agent

### Warum ?

- Coole Objekte wie PodMonitor, ServiceMonitor, PrometheusRules funktionieren
- Das ist schlecht und macht Dein unnötig schwer.
- Dann musst du nämlich die alten ScrapeConfigs verwenden (IHNNNN !)

## Prometheus / Grafana Stack installieren (advanced)

- using the kube-prometheus-stack (recommended !: includes important metrics)

### Attention: Upgrades and uninstall can be a bit tricky

- CRD's need to be deleted manually after uninstall
- Before Upgrade update the CRD's
- <https://github.com/prometheus-community/helm-charts/blob/main/charts/kube-prometheus-stack/UPGRADE.md>

### What do we want to do ?

- We want to protect prometheus with basic-auth
- We want to protect alertmanager with basic-auth
- We want to use letsencrypt

### Prerequisites

```
## 1. We have setup ingress-controller Service type:LoadBalancer -> external
## 2. We have a subdomain
## 3. Already done for you
sudo apt install apache2-utils
```

### Step 1: Create our project - folder (just to be organized)

```
cd
mkdir -p manifests
cd manifests
mkdir -p monitoring
cd monitoring
```

### Step 2: Create basic-auth

```
kubectl create ns monitoring
htpasswd -c auth admin # Enter your desired password
kubectl create secret generic prometheus-basic-auth --from-file=auth -n monitoring
```

### Step 3: Install cert-manager

```
helm repo add jetstack https://charts.jetstack.io
```

```
nano cert-manager-values.yml
```

```
crds:
  enabled: true
```

```
helm upgrade --install cert-manager jetstack/cert-manager \
  --namespace cert-manager --create-namespace --version 1.17.2 -f cert-manager-values.yml
```

### Step 4: Create ClusterIssuer

```
nano clusterissuer.yaml
```

```
apiVersion: cert-manager.io/v1
kind: ClusterIssuer
metadata:
  name: letsencrypt-prod
spec:
```

```
acme:
  email: training.tn1@t3company.de
  server: https://acme-v02.api.letsencrypt.org/directory
  privateKeySecretRef:
    name: letsencrypt-prod
  solvers:
    - http01:
        ingress:
          class: nginx
```

```
kubectl apply -f clusterissuer.yaml
```

## Step 5: Prepare Monitoring Stack (values - file)

```
nano monitoring-values.yaml
```

```
grafana:
  fullnameOverride: grafana
  enabled: true
  adminUser: admin
  adminPassword: "yourStrongPassword"
  ingress:
    enabled: true
    annotations:
      kubernetes.io/ingress.class: nginx
      cert-manager.io/cluster-issuer: letsencrypt-prod
    hosts:
      - grafana.<du>.do.t3isp.de
    path: /
    pathType: Prefix
    tls:
      - hosts:
          - grafana.<du>.do.t3isp.de
        secretName: grafana-tls

prometheus:
  fullnameOverride: prometheus
  ingress:
    enabled: true
    annotations:
      kubernetes.io/ingress.class: nginx
      nginx.ingress.kubernetes.io/auth-type: basic
      nginx.ingress.kubernetes.io/auth-secret: prometheus-basic-auth
      nginx.ingress.kubernetes.io/auth-realm: "Authentication Required"
      cert-manager.io/cluster-issuer: letsencrypt-prod
    hosts:
      - prometheus.<du>.do.t3isp.de
    paths:
      - /
    pathType: Prefix
    tls:
      - hosts:
          - prometheus.<du>.do.t3isp.de
        secretName: prometheus-tls

## Optional: Persist data
prometheusOperator:
  admissionWebhooks:
```

```

    enabled: true

alertmanager:
  fullnameOverride: alertmanager
  ingress:
    enabled: true
    annotations:
      kubernetes.io/ingress.class: nginx
      nginx.ingress.kubernetes.io/auth-type: basic
      nginx.ingress.kubernetes.io/auth-secret: prometheus-basic-auth
      nginx.ingress.kubernetes.io/auth-realm: "Authentication Required"
      cert-manager.io/cluster-issuer: letsencrypt-prod
    hosts:
      - alertmanager.<du>.do.t3isp.de
    paths:
      - /
    pathType: Prefix
    tls:
      - hosts:
          - alertmanager.<du>.do.t3isp.de
        secretName: alertmanager-tls

kube-state-metrics:
  fullnameOverride: kube-state-metrics

prometheus-node-exporter:
  fullnameOverride: node-exporter

```

## Step 6: Install with helm

```

helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm upgrade --install prometheus prometheus-community/kube-prometheus-stack -f monitoring-
values.yaml --namespace monitoring --create-namespace --version 72.3.0

```

## Step 6.5 Check, if everything works

```

kubectl -n monitoring get pods
kubectl -n cert-manager get pods

```

```

## ein neue Ressource cert-manager
## True ?
kubectl get clusterissuer
kubectl -n monitoring get certificaterequests
## Alertmanager has a problem
kubectl -n monitoring describe certificaterequests alertmanager-tls-1

kubectl -n monitoring get certificates
kubectl -n monitoring describe cert alertmanager-tls

```

## Step 7: Connect to prometheus from the outside world

```

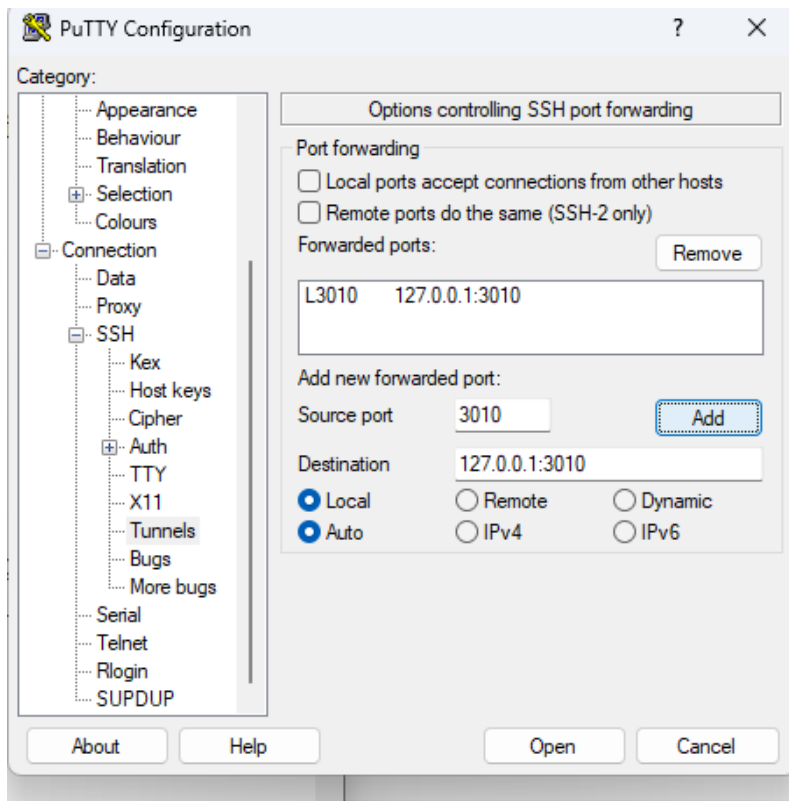
https://prometheus.<du>.do.t3isp.de

```

## Step 8: Connect to the grafana from the outside world

```
https://grafana.<du>.do.t3isp.de
```

```
## ändern in euer port + teilnehmer
## d.h. z.B. 3000 + tln1 = 3001 statt 3010
kubectl -n monitoring port-forward deploy/grafana 3010:3000 &
## if on remote - system do a ssh-tunnel
## ssh -L 3010:127.0.0.1:3010 user@remote-ip
```



## Step 9: Connect to alertmanager from the outside world

```
https://alertmanager.<du>.do.t3isp.de
```

### Attention: No persistent storage

- In this chart prometheus by default uses EmptyDir, only exists as long as pod runs
- Retention time: is 10d currently, so this long will data be there

```
prometheus-prometheus-kube-prometheus-prometheus-db:
  Type:      EmptyDir (a temporary directory that shares a pod's lifetime)
  Medium:
  SizeLimit: <unset>
```

### Set to storageclass

```
nano monitoring-values.yaml
```

```
grafana:
  fullnameOverride: grafana
  enabled: true
```



```

adminUser: admin
adminPassword: "yourStrongPassword"
ingress:
  enabled: true
  annotations:
    kubernetes.io/ingress.class: nginx
    cert-manager.io/cluster-issuer: letsencrypt-prod
  hosts:
    - grafana.<du>.do.t3isp.de
  path: /
  pathType: Prefix
  tls:
    - hosts:
        - grafana.<du>.do.t3isp.de
      secretName: grafana-tls

prometheus:
  fullnameOverride: prometheus
### That is the storageclass part
prometheusSpec:
  storageSpec:
    volumeClaimTemplate:
      spec:
        accessModes: ["ReadWriteOnce"]
        resources:
          requests:
            storage: 20Gi
        storageClassName: "standard"
#####
ingress:
  enabled: true
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/auth-type: basic
    nginx.ingress.kubernetes.io/auth-secret: prometheus-basic-auth
    nginx.ingress.kubernetes.io/auth-realm: "Authentication Required"
    cert-manager.io/cluster-issuer: letsencrypt-prod
  hosts:
    - prometheus.<du>.do.t3isp.de
  paths:
    - /
  pathType: Prefix
  tls:
    - hosts:
        - prometheus.<du>.do.t3isp.de
      secretName: prometheus-tls

## Optional: Persist data
prometheusOperator:
  admissionWebhooks:
    enabled: true

alertmanager:
  fullnameOverride: alertmanager
ingress:
  enabled: true
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/auth-type: basic

```

```

    nginx.ingress.kubernetes.io/auth-secret: prometheus-basic-auth
    nginx.ingress.kubernetes.io/auth-realm: "Authentication Required"
    cert-manager.io/cluster-issuer: letsencrypt-prod
  hosts:
    - alertmanager.<du>.do.t3isp.de
  paths:
    - /
  pathType: Prefix
  tls:
    - hosts:
        - alertmanager.<du>.t3isp.de
      secretName: alertmanager-tls

kube-state-metrics:
  fullnameOverride: kube-state-metrics

prometheus-node-exporter:
  fullnameOverride: node-exporter

```

```

## ausrollen
helm upgrade --install prometheus prometheus-community/kube-prometheus-stack -f monitoring-
values.yaml --namespace monitoring --create-namespace --version 72.3.0

```

## References:

- <https://github.com/prometheus-community/helm-charts/blob/main/charts/kube-prometheus-stack/README.md>
- <https://artifacthub.io/packages/helm/prometheus-community/prometheus>

## Prometheus / blackbox exporter

### Prerequisites

- prometheus setup with helm

### Step 1: Setup

```

helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm install my-prometheus-blackbox-exporter prometheus-community/prometheus-blackbox-exporter --
version 8.17.0 --namespace monitoring --create-namespace

```

### Step 2: Find SVC

```
kubectl -n monitoring get svc | grep blackbox
```

```
my-prometheus-blackbox-exporter    ClusterIP    10.245.183.66    <none>    9115/TCP
```

### Step 3: Test with Curl

```
kubectl run -it --rm curltest --image=curlimages/curl -- sh
```

```

## Testen nach google in shell von curl
curl http://my-prometheus-blackbox-exporter.monitoring:9115/probe?
target=google.com&module=http_2xx

```

```

## Looking for metric
probe_http_status_code 200

```

#### Step 4: Test apple-service with Curl

```
## From within curlimages/curl pod
curl http://my-prometheus-blackbox-exporter.monitoring:9115/probe?target=apple-
service.app&module=http_2xx
```

#### Step 5: Scrape Config (We want to get all services being labeled example.io/should\_be\_probed = true)

```
prometheus:
  prometheusSpec:
    additionalScrapeConfigs:
      - job_name: "blackbox-microservices"
        metrics_path: /probe
        params:
          module: [http_2xx]
        # Autodiscovery through kube-api-server
        #
https://prometheus.io/docs/prometheus/latest/configuration/configuration/#kubernetes_sd_config
        kubernetes_sd_configs:
          - role: service
        relabel_configs:
          # Example relabel to probe only some services that have "example.io/should_be_probed =
true" annotation
          - source_labels: [__meta_kubernetes_service_annotation_example_io_should_be_probed]
            action: keep
            regex: true
          - source_labels: [__address__]
            target_label: __param_target
          - target_label: __address__
            replacement: my-prometheus-blackbox-exporter:9115
          - source_labels: [__param_target]
            target_label: instance
          - action: labelmap
            regex: __meta_kubernetes_service_label_(.+)
          - source_labels: [__meta_kubernetes_namespace]
            target_label: app
          - source_labels: [__meta_kubernetes_service_name]
            target_label: kubernetes_service_name
```

#### Step 6: Test with relabeler

- <https://relabeler.promlabs.com>

#### Step 7: Scrapeconfig einbauen

```
## von kube-prometheus-grafana in values und upgraden
helm upgrade prometheus prometheus-community/kube-prometheus-stack -f values.yml --namespace
monitoring --create-namespace --version 61.3.1
```

#### Step 8: annotation in service einfügen

```
kind: Service
apiVersion: v1
metadata:
  name: apple-service
```

```

annotations:
  example.io/should_be_probed: "true"

spec:
  selector:
    app: apple
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5678 # Default port for image

```

```
kubectl apply -f service.yml
```

### Step 9: Look into Status -> Discovery Services and wait

- blackbox services should now appear under blackbox\_microservices
- and not being dropped

### Step 10: Unter <http://64.227.125.201:30090/targets?search=> gucken

- .. ob das funktioniert

### Step 11: Hauptseite (status code 200)

- Metrik angekommen `?
- [http://64.227.125.201:30090/graph?g0.expr=probe\\_http\\_status\\_code&g0.tab=1&g0.display\\_mode=lines&g0.show\\_exemplars=0&g0.range\\_input=1h](http://64.227.125.201:30090/graph?g0.expr=probe_http_status_code&g0.tab=1&g0.display_mode=lines&g0.show_exemplars=0&g0.range_input=1h)

### Step 12: pod vom service stoppen

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: apple-deployment
spec:
  selector:
    matchLabels:
      app: apple
  replicas: 8
  template:
    metadata:
      labels:
        app: apple
    spec:
      containers:
        - name: apple-app
          image: hashicorp/http-echo
          args:
            - "-text=apple-<dein-name>"

```

```
kubectl apply -f apple.yml # (deployment)
```

### Step 13: status\_code 0

- Metrik angekommen `?
- [http://64.227.125.201:30090/graph?g0.expr=probe\\_http\\_status\\_code&g0.tab=1&g0.display\\_mode=lines&g0.show\\_exemplars=0&g0.range\\_input=1h](http://64.227.125.201:30090/graph?g0.expr=probe_http_status_code&g0.tab=1&g0.display_mode=lines&g0.show_exemplars=0&g0.range_input=1h)

## Prometheus Praxis

## Nginx mit ServiceMonitor und export konfigurieren (sidecar)

### Voraussetzung:

- kube-prometheus-stack muss installiert sein -> [Kube-Prometheus-Stack installieren](#)

### Vorbereitung: Verzeichnisstruktur anlegen

```
cd ~
mkdir -p manifests
cd manifests
mkdir svcm-nginx
cd svcm-nginx
```

Alle YAML-Dateien werden in diesem Verzeichnis erstellt und mit `kubectl apply -f .` angewendet.

### 1. Namespace

```
nano 01-namespace.yaml
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: web-demo
```

```
kubectl apply -f .
```

### 2. ConfigMap: Stub Status aktivieren

```
nano 02-nginx-stubstatus-configmap.yaml
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-stubstatus
  namespace: web-demo
data:
  default.conf: |
    server {
      listen 80;
      location / {
        root /usr/share/nginx/html;
        index index.html index.htm;
      }

      location /stub_status {
        stub_status;
        allow 127.0.0.1;
        deny all;
      }
    }
  }
```

```
kubectl apply -f .
```

---

### 3. Deployment mit Sidecar (Exporter)

```
nano 03-nginx-deployment-metrics.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  namespace: web-demo
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:stable
          ports:
            - containerPort: 80
          volumeMounts:
            - name: nginx-conf
              mountPath: /etc/nginx/conf.d/default.conf
              subPath: default.conf
        - name: exporter
          image: nginx/nginx-prometheus-exporter:latest
          args:
            - "-nginx.scrape-uri=http://localhost:80/stub_status"
          ports:
            - containerPort: 9113
      volumes:
        - name: nginx-conf
          configMap:
            name: nginx-stubstatus
```

```
kubectl apply -f .
```

---

### 4. Service mit zusätzlichem Metrics-Port

```
nano 04-nginx-service-with-metrics.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  namespace: web-demo
  labels:
    app: nginx
spec:
```

```
selector:
  app: nginx
ports:
- name: http
  port: 80
  targetPort: 80
- name: metrics
  port: 9113
  targetPort: 9113
```

```
kubectl apply -f .
```

## 5. Verbindung testen

```
kubectl run -it --rm podtest --image=busybox
```

```
## in der bash
wget -O - http://nginx.web-demo:9113/metrics
exit
```

## 6. Ingress (Optional)

```
nano 06-nginx-ingress.yaml
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: nginx
  namespace: web-demo
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - host: app.tln1.do.t3isp.de
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: nginx
            port:
              number: 80
```

```
kubectl apply -f .
```

## 7. ServiceMonitor

```
nano 05-nginx-servicemonitor.yaml
```

```

apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  name: nginx
  namespace: web-demo
  labels:
    release: prometheus # muss zu Helm-Werten passen!
spec:
  selector:
    matchLabels:
      app: nginx
  namespaceSelector:
    matchNames:
      - web-demo
  endpoints:
    - port: metrics
      path: /metrics
      interval: 15s

```

```
kubectl apply -f .
```

```

## Welches Label prometheus hat, könnt ihr prüfen
kubectl -n monitoring get pods -l release=prometheus

```

```

## Ist der ServiceMonitor konfiguriert ?
kubectl -n web-demo get servicemonitor nginx
kubectl -n web-demo get smon nginx
kubectl -n web-demo describe smon nginx

```

## 8. Targets finden (in web gui)

```

## im Browser Öffnen und nach web-demon suchen
https://prometheus.<du>.do.t3isp.de/targets

## Dann menü links oben ausklappen, ganz runter scrollen
## serviceMonitor/web-demo/nginx/0
## oder
https://prometheus.tln10.do.t3isp.de/targets?pool=serviceMonitor%2Fweb-demo%2Fnginx%2F0

```

## 9. mit promql abfragen

```

1. Zunächst finden wir heraus, welche labels diese pods haben (siehe Punkt 8)
das sieht nach job="nginx" aus

Jeder ServiceMonitor (z.B. unser, der nginx heisst), wird beim Scrapen als job="
<serviceMonitorName>"
automatisch von Kubernetes abgefragt.

```

d.h. wir können Fragen

```

## (gilt dann für alle pods)
up
## gilt für alle pods in einen für bestimmten job
up {job="nginx"}
## gilt für alle pods in einem bestimmten namespace
up {namespace="web-demo"}

```



```
## and combining all endpoints with job=nginx in the namespace web-demo
up {job="nginx", namespace="web-demo"}
```

```
## Regular Expressions also work:
up{namespace="web-demo", pod=~"nginx-.*"}
```

```
##Practical - on: https://prometheus1.tln1.do.t3isp.de/query
+ enter:
up {job="nginx"} + Press "Execute"
```

>\_ up{job="nginx"}

Execute

Table

Graph

Explain

< Evaluation time >

Load time: 70ms Result series: 3


up{container="exporter", endpoint="metrics", instance="192.168.197.5:9113", job="nginx", namespace="web-demo", pod="nginx-69547b6f65-j5c8f", service="nginx"}	1
up{container="exporter", endpoint="metrics", instance="192.168.46.3:9113", job="nginx", namespace="web-demo", pod="nginx-69547b6f65-1ttq6", service="nginx"}	1
up{container="exporter", endpoint="metrics", instance="192.168.228.68:9113", job="nginx", namespace="web-demo", pod="nginx-69547b6f65-p4hj", service="nginx"}	1

You can not also click on Graph

## 10. In Grafana ein Dashboard erstellen

### Step 1: New Dashboard

Oben rechts auf neues Dashboard erstellen klicken  
-->



### Step 2: Add Visualization

# Start your new dashboard by adding a visualization

Select a data source and then query and visualize your data with charts, stats and tables or create lists, markdowns and other widgets.

+ Add visualization

Step 3: Datasource -> Prometheus (Default) auswählen / Visualisation + Query definieren

- DataSource Prometheus is bereits vorkonfiguriert

**Panel options**

**Title**

Service Availability

**Description**

**Transparent background**

☐

**Panel links**

**Repeat options**


**Tooltip**

**Tooltip mode**

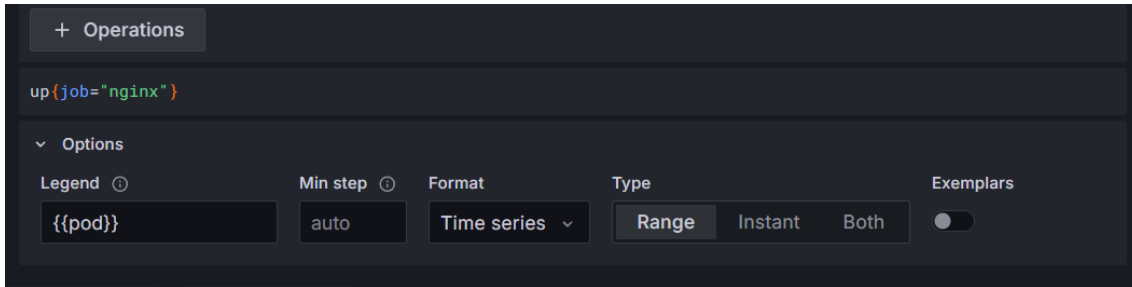
Single All Hidden

- Choose Visualisation ( Stat, Gauge, or Bar Gauge )
- Set the query: up | job | nginx

- run query



- Damit immer der pod angezeigt wird, trage dies als custom label unter der query ein



Step 4: Save Dashboard (Button oben rechts)

## Grafana - Dashboards

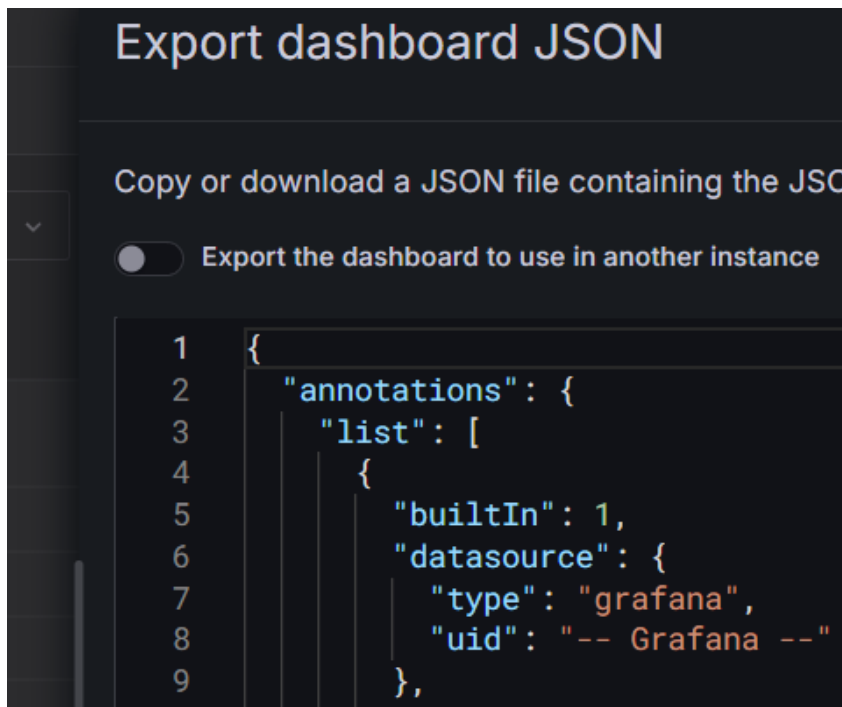
### Bestehendes Dashboard anpassen

#### Schritt 1: bestehendes Dashboard (provisioned) - exportieren bzw. importieren


- Dashboard (Node Exporter -> Nodes) aufrufen
- Oben rechts: Export -> Export as JSON

#### Schritt 2: Neues Dashboard durch import erstellen

1. Oben rechts: + Zeichen: Import Dashboard



2. Copy to clipboard
3. • und import dashboard

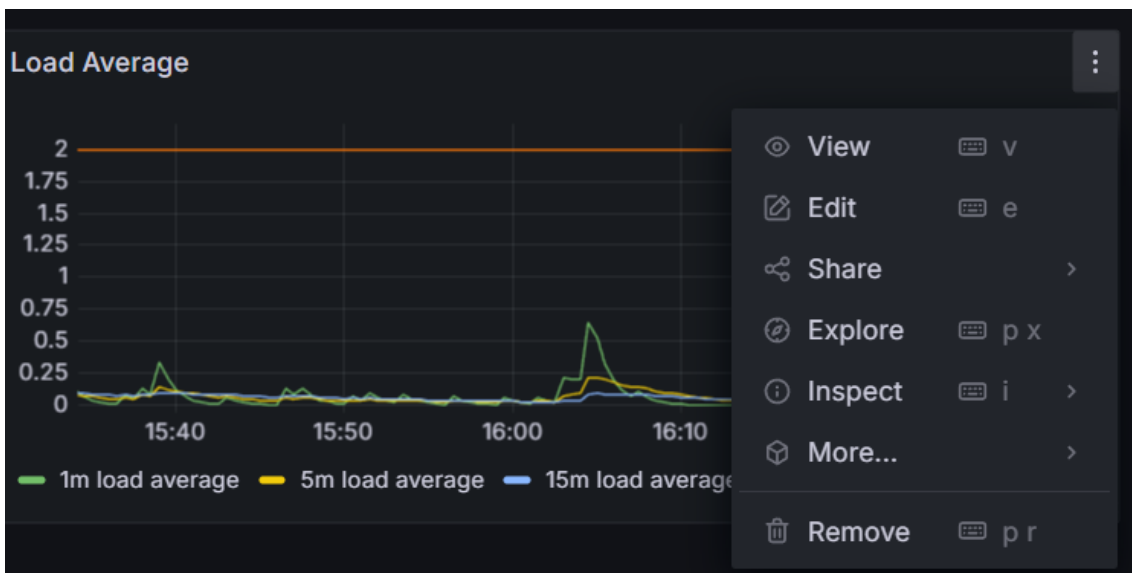


### Schritt 3: Dashboard speichern

- Achtung: Beim Speichern des Dashboards anderen Namen und andere ID angeben (einfach die uid um +1 hochzählen), vorher: change uid anklicken
- dann **Import** - Button klicken

### Schritt 4: Dashboard ausdünnen (alle nicht benötigte Panels raus)

- Wir löschen alles ausser das CPU - Panel
1. Oben rechts neben **Settings** auf **Edit** klicken
  2. Neben dem jeweiligen Panel auf die DREI-PUNKTE klicken:



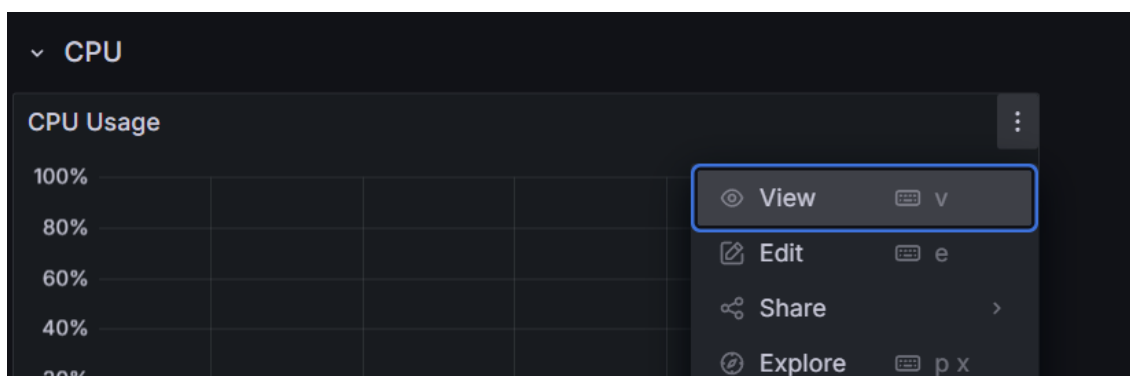
1. Auf Remove klicken
2. Alle nicht benötigten Rows löschen (Neben der Row auf den Papierkorb)

### Schritt 5: Visualisierung von CPU ändern auf Graph

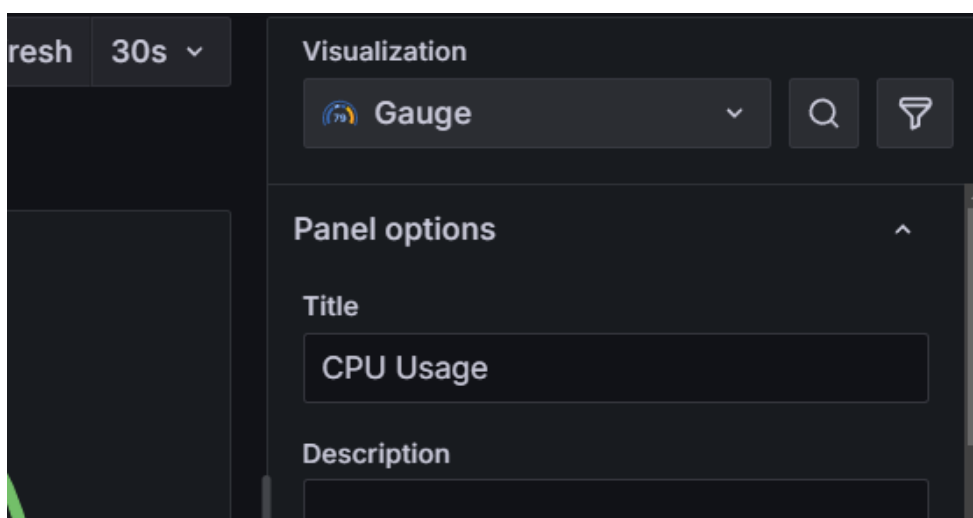
- Vorher:




- dann: rechts auf die 3 Punkte oben rechts edit



- Dann Visualisierung ändern in ->



- Nachher:




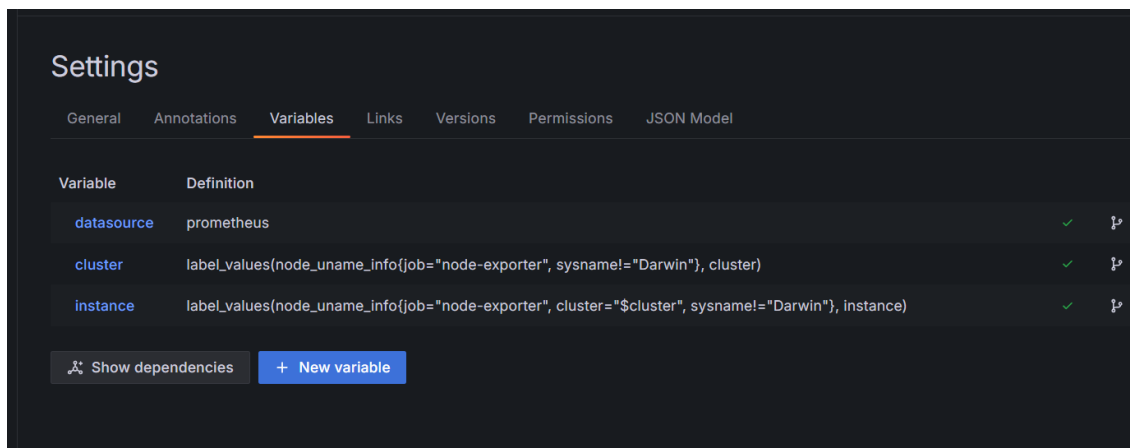
#### Schritt 6: Titel auf dynamisch ändern

- Wir setzen hier die Variable \$instance ein

The image shows a 'Panel options' configuration window. The 'Title' field is highlighted with a blue border and contains the text 'CPU Usage - \$instance'. The 'Description' field is empty. The 'Transparent background' toggle is turned off.

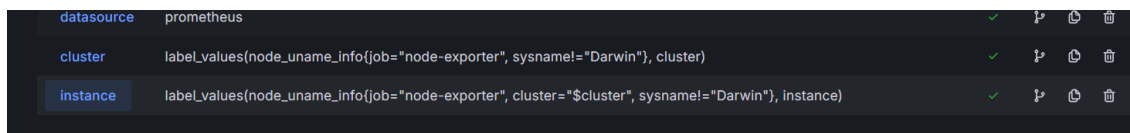
- Diese wurde bereits unter Settings -> Variablen definiert.





### Schritt 7: Variable "instance" auf Multi-Auswahl ändern

1. Unter settings: siehe 6.
2. instance anklicken



3. Multi-Value anklicken:

## Query

Query type



Label values



Label \*



instance



Metric



node\_uname\_info



Label filters



job



node-exporter



cluster

## Regex

Optional, if you want to extract part of a series name or metric node segment. Named capture groups can be used to separate the display text and value ([see examples](#)).

```
/.*(-(?<text>.*)-(?<value>.*)-.*/
```

## Sort

How to sort the values of this variable

Disabled



## Refresh

When to update the values of this variable

On dashboard load

On time range change

## Selection options



Multi-value

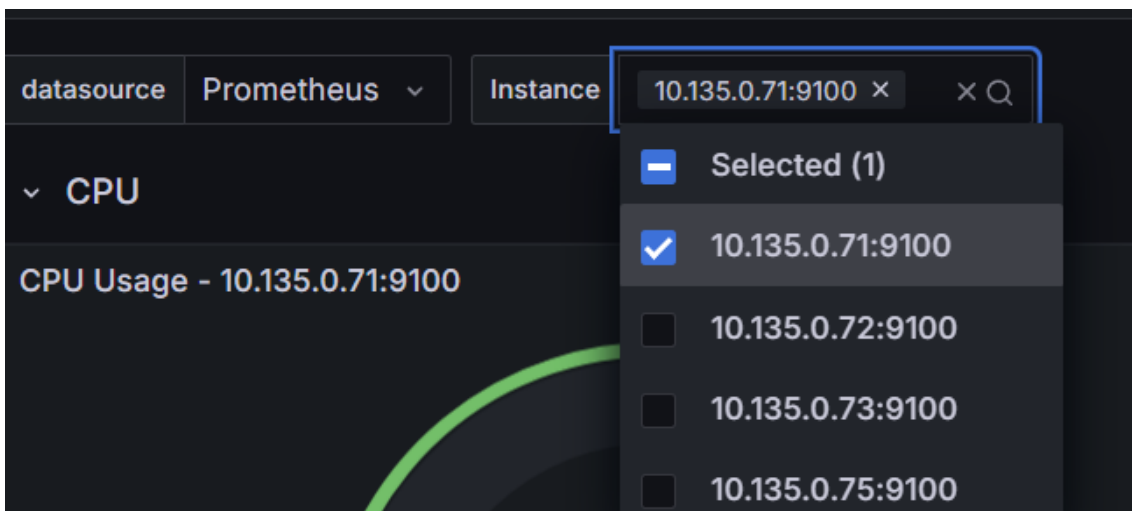

Enables multiple values to be selected at the same time

### 4. Wichtig: Dashboard speichern

Durch Multi-Value:

Dadurch lässt sich nicht nur eine Instance (ein Server), sondern mehrere auswählen





##### 5. Problem - no data begeben

- Durch umstellung auf Multi-Value muss die Query geändert werden.
- Es darf nicht mehr explizit nach einer instance gefragt werden, sondern mit regex

```
## Vorher
(
  (1 - sum without (mode) (rate(node_cpu_seconds_total{job="node-exporter",
mode=~"idle|iowait|steal", instance="$instance", cluster="$cluster"}[ $__rate_interval ])))
/ ignoring(cpu) group_left
  count without (cpu, mode) (node_cpu_seconds_total{job="node-exporter", mode="idle",
instance="$instance", cluster="$cluster"})
)
```

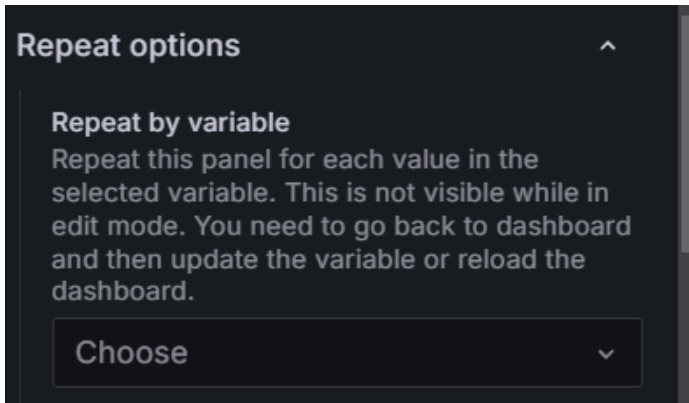
```
## Ändern in
(
  (1 - sum without (mode) (rate(node_cpu_seconds_total{job="node-exporter",
mode=~"idle|iowait|steal", instance=~"$instance", cluster="$cluster"}[ $__rate_interval ])))
/ ignoring(cpu) group_left
  count without (cpu, mode) (node_cpu_seconds_total{job="node-exporter", mode="idle",
```

```
instance=~"$instance", cluster="$cluster"}}  
)
```

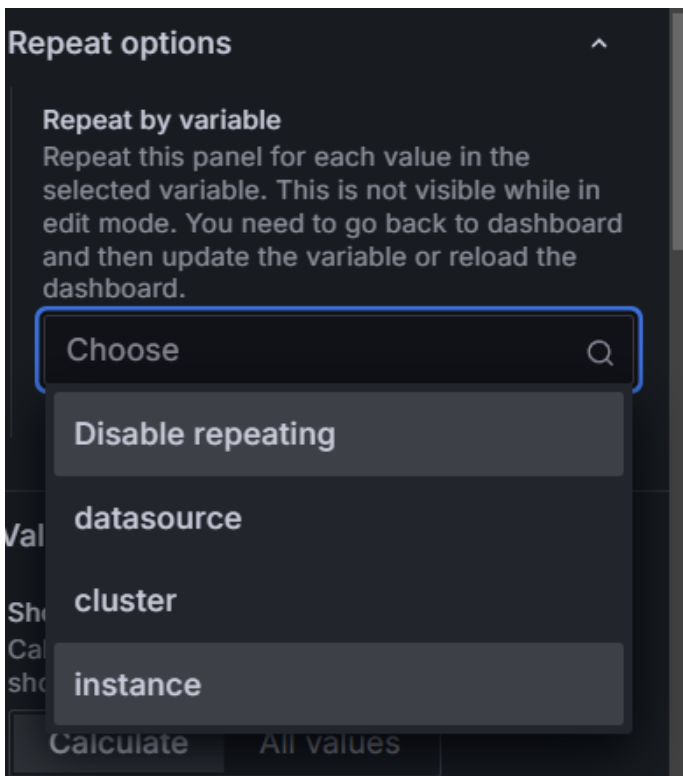
- Dashboard speichern

### Schritt 5: Auf panel repetitions umstellen

1. Bei den Panel Settings runterscrollen



2. ... und instance auswählen



3. Max per row: auf 2 stellen
4. Dashboard speichern

### Schritt 6: Testen: Dashboard muss nochmal neu geladen zu werden

1. Auf Back to Dashboard klicken

2. Die neue Ausgabe sollte erscheinen (evtl. oben bei instances nochmal alle auswählen)

## Pod und Container Dashboard

### Grafana GUI-Anleitung: Pod & Container Dashboard

#### 1. Dashboard erstellen

1. In Grafana links auf **“Dashboards”** > **“New”** > **“New Dashboard”** klicken.
2. Auf **“Add a new panel”** klicken.

---

#### 2. Panel 1: Nicht-laufende Pods pro Namespace

- **Titel:** „Fehlgeschlagene Pods je Namespace“
- **Abfrage (PromQL):**

```
count by (namespace) (kube_pod_status_phase{phase=~"Pending|Failed|Unknown"})
```

- **Panel-Typ:** Bar chart
- **X-Achse:** namespace
- **Y-Achse:** Anzahl Pods
- Speichern mit **“Apply”**

---

#### 3. Panel 2: Top 5 Container mit höchsten Restarts

- Neues Panel > Titel: „Top 5 Container Restarts“
- **Abfrage:**

```
topk(5, sum by (pod, container) (kube_pod_container_status_restarts_total))
```

- **Panel-Typ:** Table
- Optional: Transformationen aktivieren für bessere Darstellung
- Apply

---

#### 4. Panel 3: CPU-Nutzung pro Container

- Neues Panel > Titel: „CPU pro Container (millicores)“
- **Abfrage:**

```
sum by (namespace, pod, container) (
  rate(container_cpu_usage_seconds_total{container!=""}[5m])
) * 1000
```

- Panel-Typ: **Time series**
- Legende anzeigen: {{namespace}} / {{pod}} / {{container}}
- Apply

---

#### 5. Panel 4: Memory-Nutzung pro Container

- Titel: „Speicher pro Container (MiB)“
- **Abfrage:**

```
sum by (namespace, pod, container) (
  container_memory_usage_bytes{container!=""}
) / (1024 * 1024)
```

- Panel-Typ: **Time series**
- Legende anzeigen: `{{namespace}} / {{pod}} / {{container}}`
- Einheit: **MiB**
- Apply

---

## 6. Panel 5: Container Ready Status als Gauge

- Titel: „Container bereit (Ready)“
- Abfrage:

```
max by (pod, container) (kube_pod_container_status_ready)
```

- Panel-Typ: **Gauge**
- Thresholds:
  - Min: `0` , Max: `1`
  - Thresholds definieren: `0 = rot` , `1 = grün`
    - `0.5` -> rot
    - `0.99` -> gelb
    - `1` -> grün
- Einheit: none
- Apply

- 
- Dashboard bearbeiten (Zahnrad oben rechts) > **Variables** > „New“:
    - **Name:** `namespace`
    - **Type:** Query
    - **Datasource:** Prometheus
    - **Query:** `label_values(kube_pod_info, namespace)`
    - Gebe eine Liste von Werte (und zwar namespace) zurück aus der Metric `kube_pod_info`, bei der `job = kube-state-metric` ist

## Query options

Data source

Prometheus

### Query

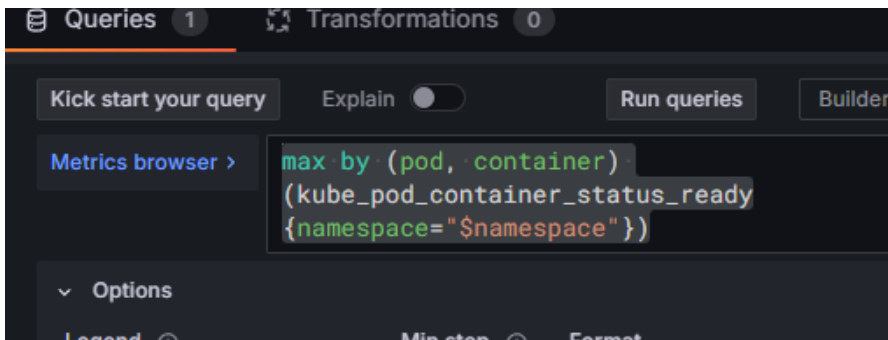
Query type	Label values
Label *	namespace
Metric	kube_pod_info
Label filters	job = kube-state-metrics

**Regex**  
Optional, if you want to extract part of a series name or metric node segment. Named capture groups can be used to separate the display text and value ([see examples](#)).

/.\*(-(?<text>.\*)-(?<value>.\*)-.\*/

- Jetzt auch für die Query anpassen (Variable verwenden)

```
max by (pod, container) (kube_pod_container_status_ready{namespace="$namespace"})
```



## 8. Save & teilen

- Dashboard speichern (Diskette-Symbol oben)
- Optional: In Ordner verschieben oder als Start-Dashboard setzen

Möchtest du eines der Panels genauer als Screenshot oder mit JSON-Code sehen? Ich kann dir auch ein Demo-Dashboard JSON liefern.

## Grafana - Alerting and Notifications

### Grafana neuen alert anlegen


#### Voraussetzung

- ServiceMonitor eingerichtet: [Nginx mit ServiceMonitor anlegen](#)

#### Basics

- am besten im über das jeweilige Panel im Dashboard (more -> new alert rule)
- Vorteil, die Query wird schon direkt übernommen

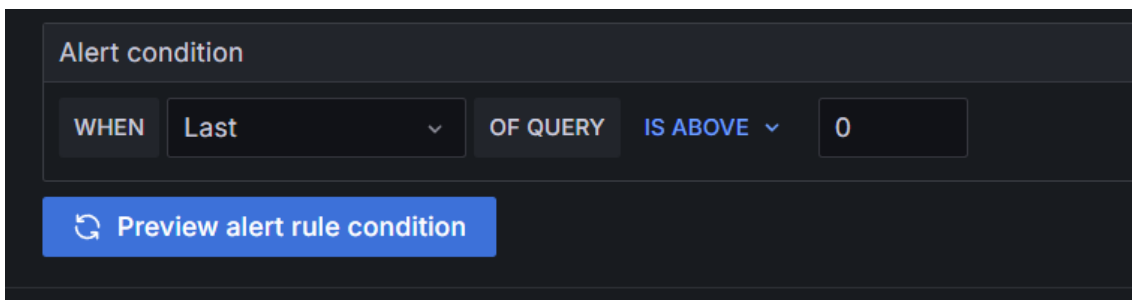
### Schritt 1: Neues Alert - Formular aufrufen



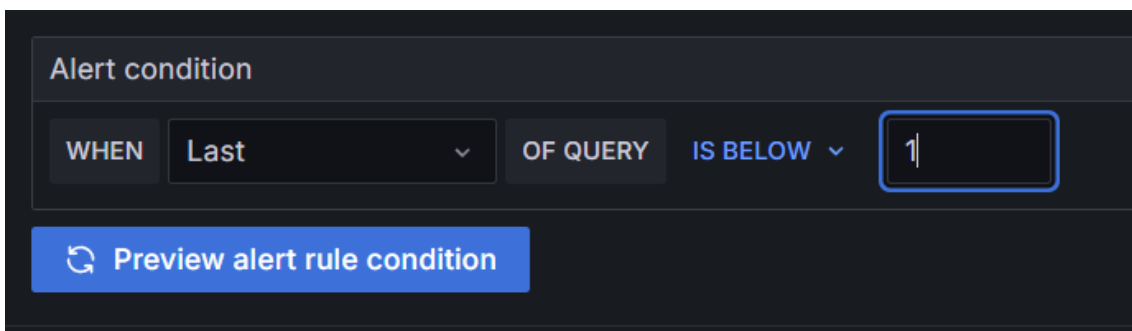
- Es wird immer zu einem Unified Alert

### Schritt 2: Wichtig: Die Alert-Condition einstellen

- Diese ist aktuell falsch



- So ist es richtig



### Schritt 3: Preview alert rule condition

- Button klicken und gucken, wie es reagiert

Alert condition

WHEN Last OF QUERY IS BELOW 1

{\_\_name\_\_="up", container="exporter", endpoint="metrics", insta

{\_\_name\_\_="up", container="exporter", endpoint="metrics", insta

{\_\_name\_\_="up", container="exporter", endpoint="metrics", insta

Preview alert rule condition

#### Schritt 4: Ein oder mehrere Labels setzen und Folder erstellen

- Wir erstellen einen neuen Folder: app1
- Labels sind u.a. wichtig für die Benachrichtigung (Die erfolgen in Form von labels)
- Wir nehmen hier

```
## label
team -> saas
```

#### Schritt 6: Set evaluation behaviour

- Wie oft soll er überprüfen ? (das kann man nach Umgebung machen, z.B)

## New evaluation group

×

Create a new evaluation group to use for this alert rule.

**Evaluation group name**  
A group evaluates all its rules over the same evaluation interval.

k8s-prod

**Evaluation interval**  
How often all rules in the group are evaluated.

1m

10s

30s

1m

5m

10m

15m

30m

1h

Cancel

Create

### Schritt 7: Configure notifications

#### Contact Points

- Contact Point auswählen
- In unserem Beispiel nehmen wir Slack

### Schritt 8: Save rule and exit

- Button oben rechts klicken

Save rule and exit

Cancel

### Schritt 9: Wo ist der Alert ?

- alert rules -> app1 -> k8s-pod



## Alert rules

Rules that determine whether an alert will fire

Search by data sources ⓘ  
All data sources

Dashboard: Select dashboard  
State: Firing Normal Pending Recovering

Rule type: Alert Recording  
Health: Ok No Data Error  
Contact point: Choose

Search ⓘ  
Q Search

View as: Grouped List State

236 rules 7 firing 143 normal 86 recording

Grafana-managed Export rules

app1 > k8s-pod 1 normal 1m ⓘ ✎

State	Name	Health	Summary	Next evaluation	Actions
Normal	New panel	ok		in a f seco	View Edit More

### Schritt 10: Testen:

- deployment löschen und im interface nachschauen

```
cd
cd manifests
cd svc-nginx
```

```
kubectl -n web-demo get deploy nginx
```

```
nano 03-nginx-deployment-metrics.yaml
```

```
## Readiness Check, einbauen, der nicht funktioniert
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  namespace: web-demo
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
```

```

- name: nginx
  image: nginx:stable
  ports:
  - containerPort: 80
  volumeMounts:
  - name: nginx-conf
    mountPath: /etc/nginx/conf.d/default.conf
    subPath: default.conf
  readinessProbe:
    exec:
      command:
      - /bin/false
    initialDelaySeconds: 5
    periodSeconds: 10
- name: exporter
  image: nginx/nginx-prometheus-exporter:latest
  args:
  - "-nginx.scrape-uri=http://localhost:80/stub_status"
  ports:
  - containerPort: 9113
  volumes:
  - name: nginx-conf

```

```

kubectl -n web-demo apply -f .
kubectl -n web-demo get pods

```

## Grafana absence alert konfigurieren - d.h. Service hat keine Pods mehr

### Prepare

- Delete deployment

```

kubectl -n web-demo get pods -l app=nginx
kubectl -n web-demo delete deploy nginx

```

### Setup Alert

2. Define query and alert condition

Define query and alert condition [Need help?](#)

Prometheus Options 6 hours, MD = 43200, Min. Interval = 30s

Kick start your query Explain

Run queries Builder Code

Metrics browser > absent(up{job="nginx"})

> Options Legend: {{pod}} Format: Time series Step: Type: Range

Alert condition

WHEN
Last
OF QUERY
IS ABOVE
0

{job="nginx"}
1 Firing

Preview alert rule condition

### 3. Add folder and labels

Organize your alert rule with a folder and set of labels. [Need help?](#)

## Click on

```
Preview and alert condition
```

## Safe and exit rules

```
Safe rule and exit
```

## Alert ausklappen und warten bis er feuert

1. Erst pending (dauert einen Moment)
2. Dann firing und es kommt eine Benachrichtigung per Slack

## ✓ Grafana Unified Alert Example: "No Data" for a Job

Use the `absent()` function

Grafana can alert when `absent(up{job="myjob"})` returns something — meaning no data is present.

### Step-by-step (in Grafana UI)

1. Go to Alerting → Alert Rules
2. Click "Create alert rule"
3. Set **Data source** to your **Prometheus**
4. Add a **query** like this:

```
absent(up{job="myjob"})
```

5. In **Conditions**, set:

- **WHEN:** Query (A) returns a number
- **IS ABOVE:** 0

This works because `absent()` returns `1` if the series is absent.

6. Under **Alert Details**, give it a name like: `No data for job "myjob"`

7. Configure **Contact Points**, **Labels**, etc.

## Grafana alert, >= pod aus replicaset nicht erreichbar

Hier ist ein vollständiges Beispiel für ein Kubernetes Deployment mit **absichtlich fehlschlagender Readiness-Probe**, das du nutzen kannst, um **Alerting zu testen**, wenn mindestens 3 Pods nicht "ready" sind:

### Beispiel-Deployment `unready-demo.yaml`

```
mkdir manifests
cd manifests
```

```
mkdir unready
cd unready
```

```
nano unready-demo.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: unready-demo
  namespace: default
spec:
  replicas: 3
  selector:
    matchLabels:
      app: unready-demo
  template:
    metadata:
      labels:
        app: unready-demo
    spec:
      containers:
        - name: demo
          image: busybox
          command: ["sh", "-c", "sleep 3600"]
          readinessProbe:
            exec:
              command:
                - /bin/false
            initialDelaySeconds: 5
            periodSeconds: 10
```

```
kubectl apply -f .
```

#### Erklärung:

- `sleep 3600` : Der Container läuft eigentlich stabil.
- `/bin/false` in der Readiness-Probe sorgt dafür, dass Kubernetes den Pod **niemals als „ready“** einstuft.
- Du kannst über `kube_pod_status_ready{condition="false"}` abfragen, wie viele Pods als „not ready“ gelten.

#### Test-Query für Alert in Grafana / Prometheus:

```
sum(kube_pod_status_ready{namespace="default", condition="false", pod=~"unready-demo-.*"}) >= 3
```

⚠ *Achte darauf, dass `kube-state-metrics` in deinem Cluster läuft und `kube_pod_status_ready` bereitstellt.*

#### ✔ So nutzt du es:

##### 1. YAML-Datei anwenden:

```
kubectl apply -f unready-demo.yaml
```

##### 2. Abfragen:

```
kubectl get pods -l app=unready-demo
```

Du wirst sehen, dass alle Pods im Status `0/1 Ready` sind.

### 3. Grafana-Alert konfigurieren auf Basis der obigen Query.

```
sum(kube_pod_status_ready(namespace="dein-namespace", condition="false", pod=~"dein-deployment-.*")) >= 3
```

---

Möchtest du auch ein **komplettes Beispiel für den passenden Grafana Alert** (als YAML oder UI-Anleitung)?

### Grafana Notifications/Contact points

To configure which **contact point** to use for an alert in a typical **alerting system** (like Prometheus Alertmanager, Grafana, etc.), the process generally involves **creating routing rules** or **setting labels** that determine how alerts are matched to specific contact points (such as email, Slack, PagerDuty, etc.).

Here's a general breakdown depending on the platform:

### **Grafana Alerting (Unified Alerting System)**

1. **Create Contact Points** via **Alerting > Contact Points**.
2. **Create Notification Policies** in **Alerting > Notification Policies**:
  - \* Define **routing rules** that match **labels or conditions**.
  - \* Attach contact points to each rule.
3. **In Alert Rules**, define **custom labels** that match the notification policies.

---

## Kubernetes Multi-Cluster (Types of setups including disadvantages/advantages)

### Recommended: Variant 1: prometheus agent + thanos/grafana stack

### Part 1: The agent side (on each cluster)

Awesome – here's a **complete working setup** for deploying **Prometheus in `agent` mode** with **`remote\_write`** using **Helm** in Kubernetes. We'll use the **`kube-prometheus-stack`** Helm chart, and assume you're pushing data to a **Thanos Receive** endpoint.

---

### Goal

- Deploy Prometheus in **Agent mode**
- Use **`remote\_write`** to send metrics to **Thanos Receive**
- Retain **no local TSDB**
- Scrape Kubernetes nodes (or other jobs)
- Run everything in a namespace like **`monitoring`**

---

### ### Step-by-step Helm Setup

#### #### 1. Add Helm Repo and Create Namespace

```
```bash
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo update

kubectl create namespace monitoring
```

## 2. Create `values-agent.yaml`

```
## values-agent.yaml

prometheus:
  prometheusSpec:
    enableRemoteWriteReceiver: false
    enableAdminAPI: false
    externalLabels:
      cluster: my-cluster

    # Enable Agent mode
    extraArgs:
      enable-feature: agent

    # No TSDB storage needed
    retention: 0h
    storageSpec: {}

    # Remote write to Thanos Receive
    remoteWrite:
      - url: http://thanos-receive.monitoring.svc:19291/api/v1/receive
        write_relabel_configs:
          - source_labels: [__name__]
            regex: ".*"
            action: keep

    # ✓ Add your scrape jobs here
    additionalScrapeConfigs:
      - job_name: 'kubernetes-nodes'
        kubernetes_sd_configs:
          - role: node
        relabel_configs:
          - action: labelmap
            regex: __meta_kubernetes_node_label_(.+)
          - source_labels: [__address__]
            target_label: __address__
            regex: (.*)\:10250
            replacement: ${1}:9100
            action: replace

## Optional: disable default recording/alerting rules
kube-prometheus-stack:
  defaultRules:
    create: false # default rules can be set to true, but agent does not use them

## Optional: disable other components like alertmanager, grafana, etc.
alertmanager:
  enabled: false
```

```
grafana:
  enabled: false
```

You can also use `configMap` for `additionalScrapeConfigs` if you prefer that method.

### 3. Install the chart with custom values

```
helm install prom-agent prometheus-community/kube-prometheus-stack \
  -n monitoring \
  -f values-agent.yaml
```

This will deploy:

- Prometheus running in **agent mode**
- Scraping Kubernetes nodes
- Streaming data to `remote_write` (e.g., Thanos Receive)

## Part 2: The monitoring cluster side (once)

Perfect — here's a complete Helm-based setup to deploy **Thanos components** in Kubernetes:

- ✓ **Thanos Receive** — accepts `remote_write` from Prometheus agents
- ✓ **Thanos Store Gateway** — reads from object storage (e.g. S3 or MinIO)
- ✓ **Thanos Query (Querier)** — unified query layer
- ✓ **Grafana** — dashboards that query Thanos Query

We'll also deploy **MinIO** as an S3-compatible store.

### Folder Structure (you'll get this as a zipped Helm lab)

```
thanos-lab/
├─ values/
│   ├── minio.yaml
│   ├── thanos-receive.yaml
│   ├── thanos-store.yaml
│   ├── thanos-query.yaml
│   └── grafana.yaml
└─ install.sh  # helper script to install everything
```

#### ✓ 1. MinIO (Object Storage)

**values/minio.yaml**

```
auth:
  rootUser: minioadmin
  rootPassword: minioadmin

defaultBuckets: thanos

resources:
  requests:
    memory: 256Mi
    cpu: 100m
```

**Install:**

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm install minio bitnami/minio -n monitoring --create-namespace -f values/minio.yaml
```

---

## ✓ 2. Thanos Receive

**values/thanos-receive.yaml**

```
objstoreConfig:
  type: s3
  config:
    bucket: "thanos"
    endpoint: "minio.monitoring.svc.cluster.local:9000"
    access_key: "minioadmin"
    secret_key: "minioadmin"
    insecure: true

receive:
  enabled: true
  tsdbRetention: 24h
  replicaLabel: replica
  hashPartitioning:
    replicationFactor: 1
    enableSharding: false

service:
  type: ClusterIP
```

**Install:**

```
helm repo add bitnami https://charts.bitnami.com/bitnami
helm install thanos-receive bitnami/thanos -n monitoring -f values/thanos-receive.yaml
```

---

## ✓ 3. Thanos Store Gateway

**values/thanos-store.yaml**

```
objstoreConfig:
  type: s3
  config:
    bucket: "thanos"
    endpoint: "minio.monitoring.svc.cluster.local:9000"
    access_key: "minioadmin"
    secret_key: "minioadmin"
    insecure: true

storegateway:
  enabled: true

service:
  type: ClusterIP
```

**Install:**

```
helm install thanos-store bitnami/thanos -n monitoring -f values/thanos-store.yaml
```



---

## ✓ 4. Thanos Query

values/thanos-query.yaml

```
query:
  enabled: true
  replicaLabel: replica

stores:
  - thanos-receive.monitoring.svc.cluster.local:10901
  - thanos-store.monitoring.svc.cluster.local:10901

service:
  type: ClusterIP
```

Install:

```
helm install thanos-query bitnami/thanos -n monitoring -f values/thanos-query.yaml
```

---

## ✓ 5. Grafana (with Thanos Query as datasource)

values/grafana.yaml

```
datasources:
  datasources.yaml:
    apiVersion: 1
    datasources:
      - name: Thanos
        type: prometheus
        access: proxy
        url: http://thanos-query.monitoring.svc.cluster.local:9090
        isDefault: true
        jsonData:
          timeInterval: "15s"
    adminPassword: "admin"
```

Install:

```
helm install grafana bitnami/grafana -n monitoring -f values/grafana.yaml
```

---

### One-liner install script (optional)

You can create a small script `install.sh` :

```
#!/bin/bash
NAMESPACE="monitoring"
helm install minio bitnami/minio -n $NAMESPACE --create-namespace -f values/minio.yaml
helm install thanos-receive bitnami/thanos -n $NAMESPACE -f values/thanos-receive.yaml
helm install thanos-store bitnami/thanos -n $NAMESPACE -f values/thanos-store.yaml
helm install thanos-query bitnami/thanos -n $NAMESPACE -f values/thanos-query.yaml
helm install grafana bitnami/grafana -n $NAMESPACE -f values/grafana.yaml
```

---

Would you like me to package this into a downloadable ZIP for you?

## Variant 2: Full prometheus in each cluster with thanos sidecar

Great! Here's a basic **Kubernetes manifest example** and also a **Helm-based setup** outline for deploying **Thanos components** alongside **Prometheus**, including long-term storage with **MinIO (S3-compatible)**. This setup is **lab-ready**, simple, and suitable for small-scale clusters or training.

---

### Thanos Setup with Helm (Prometheus + Thanos + MinIO)

#### Option A: Using Helm

##### 1. Add Repos

```
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo add bitnami https://charts.bitnami.com/bitnami
helm repo update
```

---

##### 2. Deploy MinIO (S3-compatible storage)

```
helm install minio bitnami/minio \
  --set accessKey.password=minioadmin \
  --set secretKey.password=minioadmin \
  --set defaultBuckets=thanos \
  --namespace monitoring --create-namespace
```

---

##### 3. Deploy Prometheus with Thanos Sidecar

Create `values-prometheus.yaml` :

```
extraContainers:
- name: thanos-sidecar
  image: quay.io/thanos/thanos:v0.34.0
  args:
  - sidecar
  - --tsdb.path=/data
  - --prometheus.url=http://localhost:9090
  - --objstore.config-file=/etc/thanos/objstore.yaml
  volumeMounts:
  - name: prometheus-data
    mountPath: /data
  - name: thanos-objstore
    mountPath: /etc/thanos

extraVolumes:
- name: thanos-objstore
  configMap:
    name: thanos-objstore

service:
  enabled: true
  type: ClusterIP
```

Create a ConfigMap for `objstore.yaml` :

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: thanos-objstore
  namespace: monitoring
```

```
data:
  objstore.yaml: |
    type: s3
    config:
      bucket: "thanos"
      endpoint: "minio.monitoring.svc.cluster.local:9000"
      access_key: "minioadmin"
      secret_key: "minioadmin"
      insecure: true
```

Install Prometheus:

```
helm install prometheus prometheus-community/prometheus \
  -f values-prometheus.yaml \
  --namespace monitoring
```

---

### Prometheus Responsibilities

- **Scrapes metrics** from your applications and Kubernetes targets.
- Stores those metrics locally (TSDB).
- **Has all the `scrape_configs`** (e.g., pods, services, kubelets, etc.).
- Runs as usual with no changes to scraping behavior.
- Now runs **with a Thanos Sidecar** container.

---

### Thanos Sidecar Responsibilities

- Reads data from the local TSDB (no scraping itself).
- Uploads blocks to object storage (e.g., MinIO, S3).
- Exposes a gRPC endpoint to Thanos Query or Store Gateway.
- Acts as a **bridge** between Prometheus and Thanos.

---

### Where Are Scrape Configs Defined?

- They are defined in:
  - `prometheus.yaml` if you're managing raw config.
  - Or via `values.yaml` in Helm under `serverFiles.prometheus.yaml.scrape_configs`.

Example in Helm:

```
serverFiles:
  prometheus.yaml:
    scrape_configs:
      - job_name: 'kubernetes-nodes'
        kubernetes_sd_configs:
          - role: node
        relabel_configs:
          - action: labelmap
            regex: __meta_kubernetes_node_label_(.+)
```

## Grafana Loki

### Installation von Grafana Loki - Single Instance - für Testing

#### Voraussetzung:

- Prometheus / Grafana Monitoring - Stack läuft bereits im namespace "monitoring"
- Prometheus ist als release "prometheus" mit helm installiert

#### Schritt 0: csi ausrollen (nfs-server muss eingerichtet sein)

```
helm repo add csi-driver-nfs https://raw.githubusercontent.com/kubernetes-csi/csi-driver-nfs/master/charts
helm install csi-driver-nfs csi-driver-nfs/csi-driver-nfs --namespace kube-system --version v4.11.0
```

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
  name: nfs-csi
provisioner: nfs.csi.k8s.io
parameters:
  server: 10.135.0.34
  share: /var/nfs
reclaimPolicy: Retain
volumeBindingMode: Immediate
```

```
kubectl apply -f .
```

### Schritt 1: Projektordner anlegen

```
cd ~
mkdir loki-single
cd loki-single
```

Damit wird dein Projekt im Home-Verzeichnis ( `~/loki-single` ) angelegt.

---

### Schritt 2: `values.yaml` erstellen

```
nano values.yaml
```

```
## Disabled for testing, otherwise cluster node needs way more than 8 GB of Memory
chunksCache:
  enabled: false

loki:
  auth_enabled: false
  commonConfig:
    replication_factor: 1
  schemaConfig:
    configs:
      - from: "2024-04-01"
        store: tsdb
        object_store: s3
        schema: v13
        index:
          prefix: loki_index_
          period: 24h
  pattern_ingester:
    enabled: true
  limits_config:
    allow_structured_metadata: true
    volume_enabled: true
  ruler:
```

```

    enable_api: true

minio:
  enabled: true

deploymentMode: SingleBinary

singleBinary:
  replicas: 1

## Zero out replica counts of other deployment modes
backend:
  replicas: 0
read:
  replicas: 0
write:
  replicas: 0

ingester:
  replicas: 0
querier:
  replicas: 0
queryFrontend:
  replicas: 0
queryScheduler:
  replicas: 0
distributor:
  replicas: 0
compactor:
  replicas: 0
indexGateway:
  replicas: 0
bloomCompactor:
  replicas: 0
bloomGateway:
  replicas: 0

```

### Schritt 3: Installieren

```

helm repo add grafana https://grafana.github.io/helm-charts
helm repo update

helm upgrade --install loki grafana/loki \
  --namespace loki --create-namespace --version 6.29.0 \
  -f values.yaml

```

### Schritt 4: promtail

```
nano promtail-values.yaml
```

```

config:
  clients:
    - url: http://loki-gateway.loki.svc.cluster.local/loki/api/v1/push

```

```
helm install promtail grafana/promtail --namespace loki -f promtail-values.yaml --create-namespace
```

## Ref:

- <https://grafana.com/docs/loki/latest/setup/install/helm/install-monolithic/>

## Datasource in Grafana bereitstellen per helm

### Voraussetzung:

- Prometheus / Grafana Monitoring - Stack läuft bereits im namespace "monitoring"
- Prometheus ist als release "prometheus" mit helm installiert

### Schritt 0: csi ausrollen (nfs-server muss eingerichtet sein)

```
helm repo add csi-driver-nfs https://raw.githubusercontent.com/kubernetes-csi/csi-driver-nfs/master/charts
helm install csi-driver-nfs csi-driver-nfs/csi-driver-nfs --namespace kube-system --version v4.11.0
```

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
  name: nfs-csi
provisioner: nfs.csi.k8s.io
parameters:
  server: 10.135.0.34
  share: /var/nfs
reclaimPolicy: Retain
volumeBindingMode: Immediate
```

```
kubectl apply -f .
```

### Schritt 1: Projektordner anlegen

```
cd ~
mkdir loki-single
cd loki-single
```

Damit wird dein Projekt im Home-Verzeichnis ( ~/loki-single ) angelegt.

---

### Schritt 2: values.yaml erstellen

```
nano values.yaml
```

```
## Disabled for testing, otherwise cluster node needs way more than 8 GB of Memory
chunksCache:
  enabled: false

loki:
  auth_enabled: false
  commonConfig:
    replication_factor: 1
  schemaConfig:
    configs:
      - from: "2024-04-01"
        store: tsdb
```

```

    object_store: s3
    schema: v13
    index:
      prefix: loki_index_
      period: 24h
    pattern_ingester:
      enabled: true
    limits_config:
      allow_structured_metadata: true
      volume_enabled: true
    ruler:
      enable_api: true

minio:
  enabled: true

deploymentMode: SingleBinary

singleBinary:
  replicas: 1

## Zero out replica counts of other deployment modes
backend:
  replicas: 0
read:
  replicas: 0
write:
  replicas: 0

ingester:
  replicas: 0
querier:
  replicas: 0
queryFrontend:
  replicas: 0
queryScheduler:
  replicas: 0
distributor:
  replicas: 0
compactor:
  replicas: 0
indexGateway:
  replicas: 0
bloomCompactor:
  replicas: 0
bloomGateway:
  replicas: 0

```

### Schritt 3: Installieren

```

helm repo add grafana https://grafana.github.io/helm-charts
helm repo update

helm upgrade --install loki grafana/loki \
  --namespace loki --create-namespace --version 6.29.0 \
  -f values.yaml

```

### Schritt 4: promtail

```
nano promtail-values.yaml
```

```
config:
  clients:
    - url: http://loki-gateway.loki.svc.cluster.local/loki/api/v1/push
```

```
helm install promtail grafana/promtail --namespace loki -f promtail-values.yaml --create-namespace
```

#### Ref:

- <https://grafana.com/docs/loki/latest/setup/install/helm/install-monolithic/>

### Wo finde ich Loki in Grafana ?

---

#### 1. Explore → Logs (Ad-hoc-Logsuche)

##### Schritte:

1. Links im Menü auf "**Explore**" klicken
2. Oben links im Dropdown die **Loki-Data Source** auswählen (z. B. `Loki` )
3. Du kannst jetzt:
  - Nach **Labels** filtern ( `{job="my-app"}` )
  - Per LogQL Abfragen wie `|= "error"` verwenden
  - Live-Logs anzeigen (unten rechts: **Live** aktivieren)

## Prometheus

### Prometheus-Metriktypen (engl. metric types)

#### Welche gibt es ?

- Counter
- Gauge
- Histogram
- Summary

In der Prometheus-Dokumentation werden sie auch explizit als metric types bezeichnet. Wenn du also über „eine Gauge“ sprichst, meinst du korrekt: „Eine Metrik vom Typ Gauge“

#### 1. Counter

Ein Counter ist ein **ständig wachsender Wert**. Er beginnt bei 0 und kann **nur steigen** (außer bei einem Reset, z. B. Pod-Restart).

##### Beispiel:

- `http_requests_total` – zählt, wie viele HTTP-Anfragen es gab
- `errors_total` – zählt Fehlermeldungen

**+ Nur hochzählend!** Für „Rate“-Abfragen ideal.

##### Typischer PromQL-Ausdruck:

```
rate(http_requests_total[5m])
```

Zeigt, wie viele Anfragen pro Sekunde in den letzten 5 Minuten kamen.

---

#### 2. Gauge

Ein Gauge ist ein **aktueller Messwert**, der **steigen und sinken** kann – wie ein Thermometer.



**Beispiel:**

- `memory_usage_bytes` – aktueller Speicherverbrauch
- `cpu_temperature` – aktuelle CPU-Temperatur

*Ideal für Zustände wie Auslastung, offene Verbindungen etc.*

**PromQL:**

```
memory_usage_bytes
```

zeigt den letzten bekannten Wert.

### 3. Histogram

Ein Histogram misst **Verteilungen von Werten**, z. B. Antwortzeiten. Es zählt, **wie viele Ereignisse in bestimmte Wertebereiche ("Buckets")** fallen.

**Beispiel:**

- `http_request_duration_seconds_bucket`

Diese Metrik ist gekoppelt mit:

- `_count` (Gesamtanzahl)
- `_sum` (Summe aller Werte)

*Sehr nützlich für Latenzen und Antwortzeitverteilungen.*

**PromQL-Beispiel (90. Perzentil über 5 Minuten):**

```
histogram_quantile(0.9, rate(http_request_duration_seconds_bucket[5m]))
```

### 4. Summary (ähnlich wie Histogram, aber clientseitig berechnet)

Ein Summary enthält direkt **Perzentile**, allerdings:

- weniger aggregierbar über mehrere Instanzen
- erzeugt mehr Metriken
- eher selten verwendet in modernen Setups

**Beispiel:**

- `http_request_duration_seconds{quantile="0.9"}`

*⚠ Für verteilte Systeme nicht gut skalierbar → lieber Histogram verwenden!*

**Zusammenfassung als Tabelle:**

Typ	Eigenschaften	Beispiel	Ideal für...
Counter	nur steigend	<code>http_requests_total</code>	Events, Fehler, Anfragen
Gauge	auf- und absteigend	<code>memory_usage_bytes</code>	Zustände, Nutzung
Histogram	Buckets + Summe/Count	<code>*_bucket</code> , <code>*_sum</code> , <code>*_count</code>	Latenzverteilung, SLA-Analyse
Summary	Clientseitige Perzentile	<code>*_quantile</code>	Einfache Latenzmetriken

## Kubernetes Multi-Cluster (using Thanos)

**Prerequisites: What is Thanos**

## What is Thanos ?

**Thanos** is an open-source **highly available, long-term storage solution for Prometheus**. It extends Prometheus by adding **global querying, deduplication, downsampling, and retention capabilities** across multiple Prometheus instances.

## Why use Thanos (or: What are the problems with Prometheus)

In simple terms:

Prometheus is great for monitoring, but it has **limitations**:

- **No built-in high availability** (HA)
- **No long-term storage** (TSDB is local and limited)
- **Difficult to query across multiple Prometheus servers**

**Thanos solves all that** by sitting on top of Prometheus.

## What are the key components of Thanos ?

**Key Components of Thanos:**

1. **Sidecar** – Sits next to each Prometheus, uploads data to object storage (e.g. S3, GCS, MinIO).
2. **Store Gateway** – Reads historical data from the object storage.
3. **Query** – Global query engine that federates multiple Prometheus instances.
4. **Compactor** – Compacts and down-samples metrics data to reduce storage usage.
5. **Ruler** – Allows global alerting & recording rules.
6. **Receiver (optional)** – Receives remote writes directly, useful in cloud-native setups.

## Benefits

**What You Get with Thanos:**

Feature	Benefit
Global Query View	One place to query multiple Prometheus
HA via Sidecars	Prometheus replicas + deduplication
Object Storage	S3/GCS/MinIO for infinite retention
Downsampling	Better performance for old data
Alerting Rules	Centralized alerting with Thanos Ruler

## Components

### Explanation

In a **Kubernetes context**, **Thanos** is typically deployed as a set of components (usually as Deployments, StatefulSets, and Services), each responsible for extending **Prometheus** to enable **long-term storage, high availability, and global querying**. Here's a breakdown of the **main Thanos components** and how they work **in Kubernetes**:

---

#### 1. Thanos Sidecar

- **Runs alongside each Prometheus instance** as a sidecar container.
- **Functions:**
  - Uploads Prometheus TSDB blocks to object storage (e.g., S3, GCS, MinIO).
  - Serves the Prometheus data to Thanos Query.
- **Deployment:** Typically as a container inside the same **Pod** as Prometheus.

---

#### 2. Thanos Query

- A **central query layer** that aggregates data from multiple Prometheus + Sidecar instances or other Thanos components (e.g., Store, Ruler).
- **Functions:**

- Provides a single PromQL query interface across multiple Prometheus data sources.
  - Used by **Grafana** as the data source for querying global metrics.
  - **Deployment:** Separate Deployment or StatefulSet with a Service.
- 

### 3. Thanos Store Gateway

- Reads **historical data** directly from the object store (S3, GCS, etc.).
  - **Functions:**
    - Makes historical data available to Thanos Query.
    - Doesn't collect or scrape; it's for **read-only** access to blocks in object storage.
  - **Deployment:** Typically a separate Deployment, often with a PersistentVolume for caching.
- 

### 4. Thanos Compactor

- Periodically **compacts, deduplicates, and downsamples** blocks in object storage.
  - **Functions:**
    - Reduces storage costs and speeds up queries.
    - Only one active instance should run at a time to avoid conflicts.
  - **Deployment:** CronJob or Deployment with single replica.
- 

### 5. Thanos Ruler

- Equivalent to Prometheus's **ruler**, but works across **Thanos data sources**.
  - **Functions:**
    - Runs alerting and recording rules on **global data**.
    - Can write rule results to object storage or remote write targets.
  - **Deployment:** Standalone component, often with object storage access.
- 

### 6. Thanos Bucket Web (optional)

- UI to **inspect the contents of the object storage bucket** used by Thanos.
  - **Functions:**
    - Helps debug or verify the TSDB blocks and compaction.
  - **Deployment:** Optional Deployment or sidecar for inspection.
- 

### Bonus: Common Setup Practices in Kubernetes

- **Object Storage:** S3 / GCS / MinIO is needed for Sidecar, Store, Compactor, Ruler.
  - **Service Discovery:** Thanos components use gRPC and DNS-based discovery (Kubernetes Services).
  - **Monitoring:** Often monitored by a separate Prometheus instance.
  - **HA:** Thanos Query and Store components can be scaled horizontally for HA.
- 

Would you like a sample Helm values file or a manifest example for deploying Thanos components in Kubernetes?

## Thanos Compactor

### Explanation

#### What does the Thanos Compactor do?

The **Compactor** is a component of Thanos that:

#### 1. Compacts smaller blocks into larger ones

- Prometheus writes data in 2-hour blocks.
- The sidecar uploads these raw blocks to object storage.
- Over time, this results in lots of small blocks.
- The Compactor merges them (e.g., 2h blocks → 10h → 48h → 2-week blocks), improving query performance.

#### 2. Downsamples old data (optional)

- Creates lower-resolution versions (e.g., 5m, 1h step sizes) to speed up long-range queries.
- Useful for dashboards or queries covering weeks/months of data.

### 3. Applies retention policies

- Deletes old blocks that exceed the configured retention period.

---

### So when do you need the Compactor?

You need the **Compactor** if:

- You're using **object storage** (S3, GCS, etc.) for long-term data
- You want efficient storage and faster queries
- You need **downsampling** or **retention control**

### On what does then compactor act

- The compactor act directly on the data uploaded to s3-storage

### How many times do I need to install compactor ?

- I will only need to install it once in the cluster (and i needs read/write access to the s3 data)

## Kubernetes Multi-Cluster (using Cortex - multi-tenant tsdb's)